

zib43yjdm

September 5, 2023

##Name - Indranil Bain ##Roll No. - 2020CSB039 ###Assignment No - 03 (Titanic)

Download Titanic Dataset(<https://www.kaggle.com/heptapod/titanic/version/1#>) and do initial pre-processing and train a Logistic Regression for the classifier.

```
[1]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[2]: BASE_PATH = '/content/drive/MyDrive/CSV Files - COLAB/train_and_test2.csv'
```

```
[3]: import pandas as pd
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
import matplotlib.pyplot as plt
```

```
[4]: dataset = pd.read_csv(BASE_PATH)
dataset
```

```
[4]:
```

	Passengerid	Age	Fare	Sex	sibsp	zero	zero.1	zero.2	zero.3	\
0	1	22.0	7.2500	0	1	0	0	0	0	
1	2	38.0	71.2833	1	1	0	0	0	0	
2	3	26.0	7.9250	1	0	0	0	0	0	
3	4	35.0	53.1000	1	1	0	0	0	0	
4	5	35.0	8.0500	0	0	0	0	0	0	
...	...	...	...	...	...	...	...	...	...	
1304	1305	28.0	8.0500	0	0	0	0	0	0	
1305	1306	39.0	108.9000	1	0	0	0	0	0	
1306	1307	38.5	7.2500	0	0	0	0	0	0	
1307	1308	28.0	8.0500	0	0	0	0	0	0	
1308	1309	28.0	22.3583	0	1	0	0	0	0	
	zero.4	...	zero.12	zero.13	zero.14	Pclass	zero.15	zero.16	\	
0	0	...	0	0	0	3	0	0		
1	0	...	0	0	0	1	0	0		

2	0	...	0	0	0	3	0	0
3	0	...	0	0	0	1	0	0
4	0	...	0	0	0	3	0	0
...	...	...	...	...	...	...	...	...
1304	0	...	0	0	0	3	0	0
1305	0	...	0	0	0	1	0	0
1306	0	...	0	0	0	3	0	0
1307	0	...	0	0	0	3	0	0
1308	0	...	0	0	0	3	0	0

	Embarked	zero.17	zero.18	Survived
0	2.0	0	0	0
1	0.0	0	0	1
2	2.0	0	0	1
3	2.0	0	0	1
4	2.0	0	0	0
...	...	...	...	...
1304	2.0	0	0	0
1305	0.0	0	0	0
1306	2.0	0	0	0
1307	2.0	0	0	0
1308	0.0	0	0	0

[1309 rows x 28 columns]

```
[5]: dataset.dropna()
```

```
[5]:
```

	Passengerid	Age	Fare	Sex	sibsp	zero	zero.1	zero.2	zero.3	\
0	1	22.0	7.2500	0	1	0	0	0	0	
1	2	38.0	71.2833	1	1	0	0	0	0	
2	3	26.0	7.9250	1	0	0	0	0	0	
3	4	35.0	53.1000	1	1	0	0	0	0	
4	5	35.0	8.0500	0	0	0	0	0	0	
...	...	...	...	...	...	...	...	...	...	...
1304	1305	28.0	8.0500	0	0	0	0	0	0	
1305	1306	39.0	108.9000	1	0	0	0	0	0	
1306	1307	38.5	7.2500	0	0	0	0	0	0	
1307	1308	28.0	8.0500	0	0	0	0	0	0	
1308	1309	28.0	22.3583	0	1	0	0	0	0	

	zero.4	...	zero.12	zero.13	zero.14	Pclass	zero.15	zero.16	\
0	0	...	0	0	0	3	0	0	
1	0	...	0	0	0	1	0	0	
2	0	...	0	0	0	3	0	0	
3	0	...	0	0	0	1	0	0	
4	0	...	0	0	0	3	0	0	
...	...	...	...	...	...	...	...	...	...

1304	0	...	0	0	0	3	0	0
1305	0	...	0	0	0	1	0	0
1306	0	...	0	0	0	3	0	0
1307	0	...	0	0	0	3	0	0
1308	0	...	0	0	0	3	0	0

	Embarked	zero.17	zero.18	2urvived
0	2.0	0	0	0
1	0.0	0	0	1
2	2.0	0	0	1
3	2.0	0	0	1
4	2.0	0	0	0
...	...	...	...	...
1304	2.0	0	0	0
1305	0.0	0	0	0
1306	2.0	0	0	0
1307	2.0	0	0	0
1308	0.0	0	0	0

[1307 rows x 28 columns]

```
[6]: dataset.columns
```

```
[6]: Index(['Passengerid', 'Age', 'Fare', 'Sex', 'sibsp', 'Parch', 'Pclass', 'Embarked', '2urvived'],
          dtype='object')
```

```
[7]: dataset = dataset[
    filter(
        lambda colName: "zero" not in colName,
        dataset.columns
    )
]
dataset = dataset.drop("Passengerid", axis=1)
dataset
```

	Age	Fare	Sex	sibsp	Parch	Pclass	Embarked	2urvived
0	22.0	7.2500	0	1	0	3	2.0	0
1	38.0	71.2833	1	1	0	1	0.0	1
2	26.0	7.9250	1	0	0	3	2.0	1
3	35.0	53.1000	1	1	0	1	2.0	1
4	35.0	8.0500	0	0	0	3	2.0	0
...	...	...	...	...	...	...	...	...
1304	28.0	8.0500	0	0	0	3	2.0	0

1305	39.0	108.9000	1	0	0	1	0.0	0
1306	38.5	7.2500	0	0	0	3	2.0	0
1307	28.0	8.0500	0	0	0	3	2.0	0
1308	28.0	22.3583	0	1	1	3	0.0	0

[1309 rows x 8 columns]

```
[14]: from sklearn.preprocessing import OneHotEncoder

def one_hot_encode(X: "pd.DataFrame", col_name: "str") -> "pd.DataFrame":
    encoder = OneHotEncoder()
    encoded_df = pd.DataFrame(
        encoder.fit_transform(X[[col_name]]).toarray(),
        index=X.index,
        columns=encoder.get_feature_names_out()
    )
    X = X.join(encoded_df)
    X = X.drop(col_name, axis=1)

    return X
```

```
[15]: columns_to_encode = ["Pclass", "Embarked", "Sex"]

for column in columns_to_encode:
    dataset = one_hot_encode(dataset, column)

dataset
```

```
[15]:
```

	Age	Fare	sibsp	Parch	Survived	Pclass_1	Pclass_2	Pclass_3	\
0	22.0	7.2500	1	0	0	0.0	0.0	1.0	
1	38.0	71.2833	1	0	1	1.0	0.0	0.0	
2	26.0	7.9250	0	0	1	0.0	0.0	1.0	
3	35.0	53.1000	1	0	1	1.0	0.0	0.0	
4	35.0	8.0500	0	0	0	0.0	0.0	1.0	
...	...	...	...	...	...	...	...	...	
1304	28.0	8.0500	0	0	0	0.0	0.0	1.0	
1305	39.0	108.9000	0	0	0	1.0	0.0	0.0	
1306	38.5	7.2500	0	0	0	0.0	0.0	1.0	
1307	28.0	8.0500	0	0	0	0.0	0.0	1.0	
1308	28.0	22.3583	1	1	0	0.0	0.0	1.0	
	Embarked_0.0	Embarked_1.0	Embarked_2.0	Embarked_nan	Sex_0	Sex_1			
0	0.0	0.0	1.0	0.0	1.0	0.0			
1	1.0	0.0	0.0	0.0	0.0	1.0			
2	0.0	0.0	1.0	0.0	0.0	1.0			
3	0.0	0.0	1.0	0.0	0.0	1.0			
4	0.0	0.0	1.0	0.0	1.0	0.0			

...	...	...	...	...	...	...
1304	0.0	0.0	1.0	0.0	1.0	0.0
1305	1.0	0.0	0.0	0.0	0.0	1.0
1306	0.0	0.0	1.0	0.0	1.0	0.0
1307	0.0	0.0	1.0	0.0	1.0	0.0
1308	1.0	0.0	0.0	0.0	1.0	0.0

[1309 rows x 14 columns]

```
[19]: # Age and Fare needs to be standardized
from sklearn.preprocessing import StandardScaler
def standardize(df: "pd.DataFrame", col_name: "str") -> "pd.DataFrame":

    scaler = StandardScaler()

    df[[col_name]] = pd.DataFrame(
        data=scaler.fit_transform(df[[col_name]]),
        index=df.index,
        columns=[col_name]
    )
    return df
```

```
[20]: columns_to_standardize = ['Age', "Fare", 'sibsp', "Parch"]

for column in columns_to_standardize:
    dataset = standardize(dataset, column)

dataset
```

```
[20]:
```

	Age	Fare	sibsp	Parch	Survived	Pclass_1	Pclass_2	\
0	-0.581628	-0.503291	0.481288	-0.445000	0	0.0	0.0	
1	0.658652	0.734744	0.481288	-0.445000	1	1.0	0.0	
2	-0.271558	-0.490240	-0.479087	-0.445000	1	0.0	0.0	
3	0.426099	0.383183	0.481288	-0.445000	1	1.0	0.0	
4	0.426099	-0.487824	-0.479087	-0.445000	0	0.0	0.0	
...	...	...	...	...	...	...	...	
1304	-0.116523	-0.487824	-0.479087	-0.445000	0	0.0	0.0	
1305	0.736169	1.462034	-0.479087	-0.445000	0	1.0	0.0	
1306	0.697411	-0.503291	-0.479087	-0.445000	0	0.0	0.0	
1307	-0.116523	-0.487824	-0.479087	-0.445000	0	0.0	0.0	
1308	-0.116523	-0.211184	0.481288	0.710763	0	0.0	0.0	

	Pclass_3	Embarked_0.0	Embarked_1.0	Embarked_2.0	Embarked_nan	Sex_0	\
0	1.0	0.0	0.0	1.0	0.0	1.0	
1	0.0	1.0	0.0	0.0	0.0	0.0	
2	1.0	0.0	0.0	1.0	0.0	0.0	
3	0.0	0.0	0.0	1.0	0.0	0.0	

4	1.0	0.0	0.0	1.0	0.0	1.0
...	...	...	...	...	...	...
1304	1.0	0.0	0.0	1.0	0.0	1.0
1305	0.0	1.0	0.0	0.0	0.0	0.0
1306	1.0	0.0	0.0	1.0	0.0	1.0
1307	1.0	0.0	0.0	1.0	0.0	1.0
1308	1.0	1.0	0.0	0.0	0.0	1.0

	Sex_1
0	0.0
1	1.0
2	1.0
3	1.0
4	0.0
...	...
1304	0.0
1305	1.0
1306	0.0
1307	0.0
1308	0.0

[1309 rows x 14 columns]

```
[22]: # Preprocessing Done, lets move to model
X = dataset.drop('Survived', axis=1)
y = dataset[['Survived']]
```

```
[23]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y)
```

```
[25]: # make, train, and score the model
from sklearn.linear_model import LogisticRegression
model = LogisticRegression().fit(X_train, y_train.iloc[:,0])
accuracy = model.score(X_test, y_test)
print(f"accuracy = {accuracy}")
```

accuracy = 0.7652439024390244

2. Analyze and control the overfitting by varying the inverse of regularization strength parameter (0.1, 0.25, 0.5, 0.75, 0.9) and plot the accuracy graph for the test set.

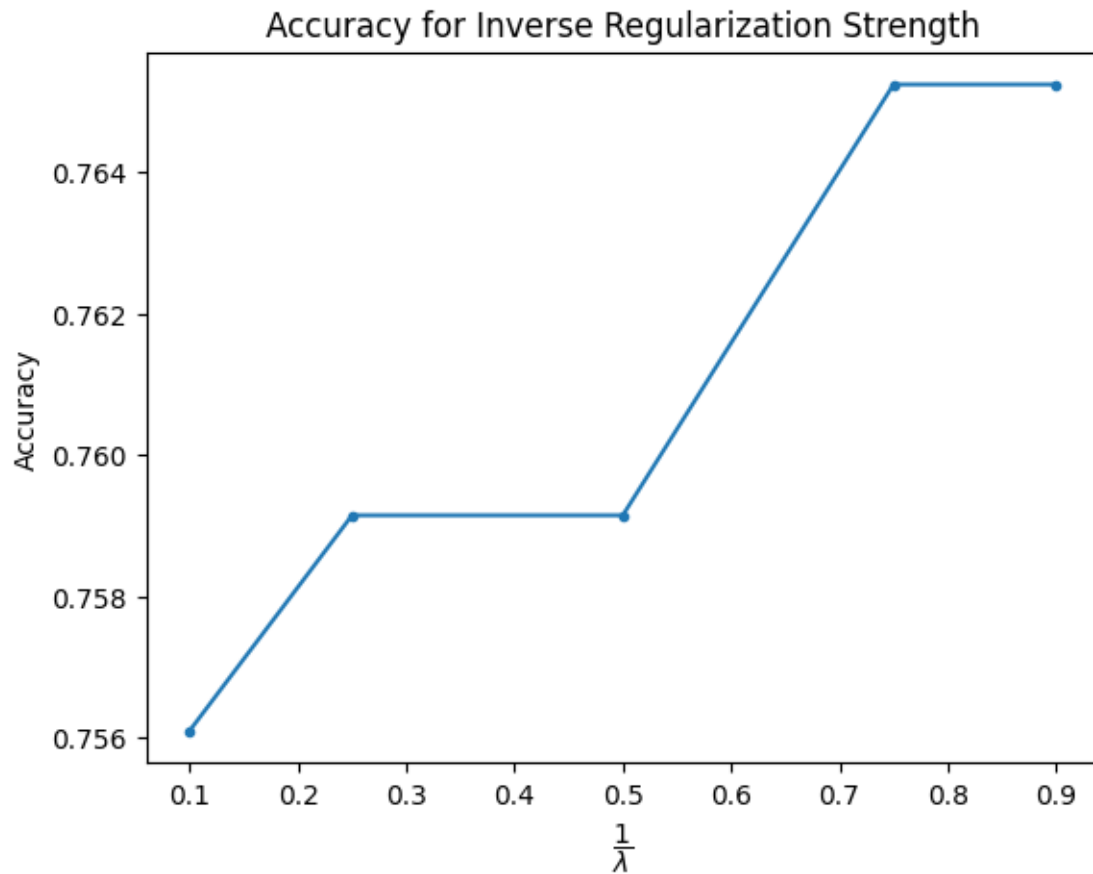
```
[26]: import matplotlib.pyplot as plt

def get_acc_log_reg( X_train: "pd.DataFrame", X_test: "pd.DataFrame", y_train:
    ↪ "pd.DataFrame", y_test: "pd.DataFrame", c=1.0
) -> "float":
    return LogisticRegression(C=c)\
        .fit(X_train, y_train.iloc[:, 0])\
```

```

        .score(X_test, y_test)
inv_reg_strs = (0.1, 0.25, 0.5, 0.75, 0.9)
accuracies = [get_acc_log_reg(X_train, X_test, y_train, y_test, c) for c in
               inv_reg_strs]
plt.plot(inv_reg_strs, accuracies, '-.')
plt.title("Accuracy for Inverse Regularization Strength")
plt.xlabel(r"$\frac{1}{\lambda}$")
plt.ylabel("Accuracy")
plt.show()

```



```

[28]: pd.DataFrame(
        data = zip(inv_reg_strs, accuracies),
        columns=['inv_reg_str', 'accuracy']
    )

```

```

[28]:
   inv_reg_str  accuracy
0         0.10  0.756098
1         0.25  0.759146
2         0.50  0.759146

```

```
3         0.75  0.765244
4         0.90  0.765244
```

3. Using the same dataset train a Decision Tree classifier and vary the maximum depth of the tree to train at least 5 classifiers to analyze the effectiveness.

```
[30]: from sklearn.tree import DecisionTreeClassifier
def get_acc_dec_tree(
    X_train: "pd.DataFrame",
    X_test: "pd.DataFrame",
    y_train: "pd.DataFrame",
    y_test: "pd.DataFrame",
    max_depth=1
) -> "float":
    return DecisionTreeClassifier(max_depth=max_depth)\
        .fit(X_train, y_train)\
        .score(X_test, y_test)
max_depths = range(1, 35)
train_accuracies = [get_acc_dec_tree(X_train, X_train, y_train, y_train, max_d)
    ↪for max_d in max_depths]
test_accuracies = [get_acc_dec_tree(X_train, X_test, y_train, y_test, max_d)
    ↪for max_d in max_depths]

plt.plot(max_depths, train_accuracies, "-.", label='Train')
plt.plot(max_depths, test_accuracies, "-.", label='Test')
plt.title("DecisionTreeClassifier Max Depth vs Accuracy")
plt.xlabel("Max Depth")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```



